

COMP 421: Files & Databases

Lecture 17: 2 Phase, 2 Locking

Announcements

Leaderboard policy: mostly, defer to original course policy, 5 bonus slots taken from "Test Only" at end of day today

Project 3 releases later today

Project 2 scores/leaderboard... next week

Last Class

Conflict Serializable

- Verify using either the “swapping” method or dependency graphs.
- Any DBMS that says that they support “serializable” isolation does this.

View Serializable

- No efficient way to verify.
- No DBMS that supports this.

Observation


We need a way to guarantee that all execution schedules are correct (i.e., serializable) without knowing the entire schedule ahead of time.

Solution: Use locks to protect database objects.

Locks vs. Latches

	<i>Locks</i>	<i>Latches</i>
Separate...	Transactions	Workers (threads, processes)
Protect...	Database Contents	In-Memory Data Structures
During...	Entire Transactions	Critical Sections
Modes...	Shared, Exclusive, Update, Intention	Read, Write
Deadlock ...by...	Detection & Resolution Waits-for, Timeout, Aborts	Avoidance Coding Discipline
Kept in...	Lock Manager	Protected Data Structure

Executing with Locks



Today's Agenda

Lock Types

Two-Phase Locking

Deadlock Detection + Prevention

Hierarchical Locking

Compatibility of lock modes

The following table shows the compatibility of any two modes for page and row locks. No question of compatibility arises between page and row locks, because a partition or table space cannot use both page and row locks.

Table 1. Compatibility matrix of page lock and row lock modes

Lock mode	Share (S-lock)	Update (U-lock)
Share (S-lock)	Yes	Yes
Update (U-lock)	Yes	No
Exclusive (X-lock)		

Compatibility for table space locks
modes for partition, table space, or

Table 2. Compatibility of table and

Lock Mode	IS	IX	S
IS	Yes	Yes	Yes
IX	Yes	Yes	No
S	Yes	No	Yes
U	Yes	No	No
SIX	Yes	No	No
X	No	No	No

Existing granted mode

Requested mode

Intent shared (IS)

IS S U

Yes Yes

Shared (S)

Yes Yes

Yes Yes

Update (U)

Yes No

Intent exclusive (IX)

Yes No

Yes No

Shared with intent exclusive (SIX)

No No



Table 13-3 Summary of Table Locks

SQL Statement	Mode of Table Lock	Lock Modes Permitted?			
		RS	RX	S	SRX
SELECT ... FROM table...	none	Y	Y	Y	Y
INSERT INTO table	RX	Y	Y	N	N
...					
UPDATE table...	RX	Y*	Y*	N	N
DELETE FROM table	RX	Y*	Y*	N	N
...					
SELECT ... FROM table FOR UPDATE OF ...	RS	Y*	Y*	Y*	Y*
LOCK TABLE table IN ROW SHARE MODE	RS	Y	Y	Y	Y
LOCK TABLE table IN ROW EXCLUSIVE MODE	RX	Y	Y	N	N
LOCK TABLE table IN SHARE MODE	S	Y	N	Y	N
LOCK TABLE table IN SHARE ROW EXCLUSIVE MODE	SRX	Y	N	N	N
LOCK TABLE table IN EXCLUSIVE MODE	X	N	N	N	N

Table-level lock type compatibility is summarized in the following matrix:

	X	IX	S	IS
X	Conflict	Conflict	Conflict	Conflict
IX	Conflict	Compatible	Conflict	Conflict
S	Conflict	Conflict	Compatible	Compatible
IS	Conflict	Compatible	Compatible	Compatible



Table 13.2. Conflicting Lock Modes

Requested Lock Mode	Existing Lock Mode						
	ACCESS	SHARE	ROW SHARE	SHARE ROW EXCL.	SHARE UPDATE EXCL.	SHARE ROW EXCL.	EXCL.
ACCESS SHARE							X
ROW SHARE			X			X	
ROW EXCL.			X			X	
SHARE UPDATE EXCL.			X			X	
SHARE			X			X	
SHARE ROW EXCL.			X			X	
EXCL.	X		X			X	
ACCESS EXCL.	X	X	X	X	X	X	X

Executing With Locks

Transactions request locks (or upgrades).


Lock manager grants or blocks requests.

Transactions release locks.

Lock manager updates its internal lock-table.

→ It keeps track of what transactions hold what locks and what transactions are waiting to acquire any locks.

Executing With Locks



Concurrency Control Protocol

Two-phase locking (2PL) is a concurrency control protocol that determines whether a txn can access an object in the database at runtime.

The protocol does not need to know all the queries that a txn will execute ahead of time.

Two-Phase Locking

Phase #1: Growing


- Each txn requests the locks that it needs from the DBMS's lock manager.
- The lock manager grants/denies lock requests.

Phase #2: Shrinking

- The txn is allowed to only release/downgrade locks that it previously acquired. It cannot acquire new locks.


Two-Phase Locking

The txn is not allowed to acquire/upgrade locks after the growing phase finishes.




Two-Phase Locking

The txn is not allowed to acquire/upgrade locks after the growing phase finishes.



Executing With 2PL




Two-Phase Locking

2PL on its own is sufficient to guarantee conflict serializability because it generates schedules whose precedence graph is acyclic.

But it is subject to cascading aborts.

2PL: Cascading Aborts

Schedule



This is a permissible schedule in 2PL, but the DBMS has to also abort T_2 when T_1 aborts.

Any information about T_1 cannot be “leaked” to the outside world.

Any computation performed
Wasted work! will be rolled back.

2PL Observations

There are potential schedules that are serializable but would not be allowed by 2PL because locking limits concurrency.

→ Most DBMSs prefer correctness before performance.

May still have “dirty reads”.

→ Solution: **Strong Strict 2PL (aka Rigorous 2PL)**


May lead to deadlocks.

→ Solution: **Detection or Prevention**

Strong Strict Two-Phase Locking

The txn is only allowed to release locks after it has ended (i.e., committed or aborted).

Allows only conflict serializable schedules, but it is often stronger than needed for some apps.



Strong Strict Two-Phase Locking

A schedule is **strict** if a value written by a txn is not read or overwritten by other txns until that txn finishes.

Advantages:

- Does not incur cascading aborts.
- Aborted txns can be undone by just restoring original values of modified tuples.

Examples

T_1 – Move \$100 from Angela's account (**A**) to Ben's account (**B**).

T_2 – Compute the total amount in all accounts and return it to the application.

T_1


```
BEGIN  
A=A-100  
B=B+100  
COMMIT
```

T_2

```
BEGIN  
ECHO A+B  
COMMIT
```

Non-2PL Example

Schedule



T ₁	T ₂
BEGIN X-LOCK(A) R(A) A=A-100 W(A) UNLOCK(A) X-LOCK(B)  R(B) B=B+100 W(B) UNLOCK(B) COMMIT	BEGIN S-LOCK(A)  R(A) UNLOCK(A) S-LOCK(B) R(B) UNLOCK(B) ECHO A+B COMMIT

Initial Database State


A=1000, B=1000

T₂ Output

A+B=1900

2PL Example

Schedule



Initial Database State

A=1000, B=1000

T₂ Output

A+B=2000

Strong Strict 2PL Example

Schedule



T ₁	T ₂
BEGIN X-LOCK(A) R(A) A=A-100 W(A) X-LOCK(B) R(B) B=B+100 W(B) UNLOCK(A) UNLOCK(B) COMMIT	BEGIN S-LOCK(A) R(A) S-LOCK(B) R(B) ECHO A+B UNLOCK(A) UNLOCK(B) COMMIT

Initial Database State

A=1000, B=1000

T₂ Output

A+B=2000

Universe of Schedules

All Schedules

View Serializable

Conflict Serializable

*No Cascading
Aborts*

Strong Strict 2PL

Serial

2PL Observations

There are potential schedules that are serializable but would not be allowed by 2PL because locking limits concurrency.

→ Most DBMSs prefer correctness before performance.


May still have “dirty reads”.

→ Solution: **Strong Strict 2PL (aka Rigorous 2PL)**

May lead to deadlocks.

→ Solution: **Detection or Prevention**

It Just Got Real



2PL Deadlocks

A **deadlock** is a cycle of transactions waiting for locks to be released by each other.

Two ways of dealing with deadlocks:

- **Approach #1: Deadlock Detection**
- **Approach #2: Deadlock Prevention**

Deadlock Detection


The DBMS creates a **waits-for** graph to keep track of what locks each txn is waiting to acquire:

- Nodes are transactions
- Edge from T_i to T_j if T_i is waiting for T_j to release a lock.


The system periodically checks for cycles in ***waits-for*** graph and then decides how to break it.

Deadlock Detection

Schedule



Waits-For Graph



Deadlock Handling

When the DBMS detects a deadlock, it will select a “victim” txn to rollback to break the cycle.

The victim txn will either restart or abort (more common) depending on how it was invoked.

There is a trade-off between the frequency of checking for deadlocks and how long txns wait before deadlocks are broken.

Deadlock Handling: Victim Selection

Selecting the proper victim depends on a lot of different variables....

- By age (lowest timestamp)
- By progress (least/most queries executed)
- By the # of items already locked
- By the # of txns that we have to rollback with it

We also should consider the # of times a txn has been restarted in the past to prevent starvation.

Deadlock Handling: Rollback Length

After selecting a victim txn to abort, the DBMS can also decide on how far to rollback the txn's changes.

Approach #1: Completely

- Rollback entire txn and tell the application it was aborted.

Approach #2: Partial (Savepoints)

- DBMS rolls back a portion of a txn (to break deadlock) and then attempts to re-execute the undone queries.

Deadlock Prevention

When a txn tries to acquire a lock that is held by another txn, the DBMS kills one of them to prevent a deadlock.

This approach does not require a ***waits-for*** graph or detection algorithm.

Deadlock Prevention

Assign priorities based on timestamps:

→ Older Timestamp = Higher Priority

Wait-Die (“Old Waits for Young”)


- If *requesting txn* has higher priority than *holding txn*,
then *requesting txn* waits for *holding txn*.
- Otherwise *requesting txn* aborts.

Wound-Wait (“Young Waits for Old”)

- If *requesting txn* has higher priority than *holding txn*,
then *holding txn* aborts and releases lock.
- Otherwise *requesting txn* waits.

Deadlock Prevention

(T_1 older than T_2)




Wait-Die

T_1 waits

Wound-Wait

T_2 aborts



Wait-Die

T_2 aborts

Wound-Wait


T_2 waits

Deadlock Prevention


Why do these schemes guarantee no deadlocks?

Timestamps define a total ordering on transactions.

Waits-For Graph



Age Graph




Deadlock Prevention

Why do these schemes guarantee no deadlocks?


Timestamps define a total ordering on transactions.

Waits-For Graph

Possible wait-die edges



Age Graph



In both schemes, arrows in "wait-for" only point in one direction


Deadlock Prevention

Why do these schemes guarantee no deadlocks?


Timestamps define a total ordering on transactions.

Waits-For Graph

Possible wound-wait edges



Age Graph



In both schemes, arrows in "wait-for" only point in one direction

Deadlock Prevention

Why do these schemes guarantee no deadlocks?

Timestamps define a total ordering on transactions.

When a txn restarts, what is its (new) priority?

Its original timestamp to prevent it from getting starved for resources like an old man at a corrupt senior center.

Observation

All these examples have a one-to-one mapping from database objects to locks.

If a txn wants to update one billion tuples, then it must acquire one billion locks.

Acquiring locks is a more expensive operation than acquiring a latch even if that lock is available.

Lock Granularities

When a txn wants to acquire a “lock”, the DBMS can decide the granularity (i.e., scope) of that lock.


→ Attribute? Tuple? Page? Table?

The DBMS should ideally obtain fewest number of locks that a txn needs.

Trade-off between parallelism versus overhead.

→ Fewer Locks, Larger Granularity vs. More Locks, Smaller Granularity.

Database Lock Hierarchy



Intention Locks

An **intention lock** allows a higher-level node to be locked in **shared** or **exclusive** mode without having to check all descendent nodes.

If a node is locked in an intention mode, then some txn is doing explicit locking at a lower level in the tree.

Intention Locks

Intention-Shared (IS)

- Indicates explicit locking at lower level with **S** locks.
- Intent to get **S** lock(s) at finer granularity.

Intention-Exclusive (IX)

- Indicates explicit locking at lower level with **X** locks.
- Intent to get **X** lock(s) at finer granularity.

Shared+Intention-Exclusive (SIX)

- The subtree rooted by that node is locked explicitly in **S** mode and explicit locking is being done at a lower level with **X** locks.

Compatibility Matrix

		T_2 Wants				
		IS	IX	S	SIX	X
T_1 Holds	IS	✓	✓	✓	✓	✗
	IX	✓	✓	✗	✗	✗
	S	✓	✗	✓	✗	✗
	SIX	✓	✗	✗	✗	✗
	X	✗	✗	✗	✗	✗

Locking Protocol

Each txn obtains appropriate lock at highest level of the database hierarchy.

To get **S** or **IS** lock on a node, the txn must hold at least **IS** on parent node.

To get **X**, **IX**, or **SIX** on a node, must hold at least **IX** on parent node.

Example

T_1 – Get the balance of Angela's bank account.

T_2 – Increase Ben's account balance by 1%.


What locks should these txns obtain?

- Exclusive + Shared for leaf nodes of lock tree.
- Special Intention locks for higher levels.

Example – Two-level Hierarchy

Read Angela's record in R.


Update Ben's record in R.




Example – Three Txns

Assume three txns execute at same time:

- T_1 – Scan all tuples in R and update one tuple.
- T_2 – Read a single tuple in R .
- T_3 – Scan all tuples in R .



Example – Three Txns



Lock Escalation

The DBMS can automatically switch to coarser-grained locks when a txn acquires too many low-level locks.

This reduces the number of requests that the lock manager must process.

Locking In Practice

Applications typically do not acquire a txn's locks manually (i.e., explicit SQL commands).

Sometimes you need to provide the DBMS with hints to help it to improve concurrency.

- Update a tuple after reading it.
- Skip any tuple that is locked.

Explicit locks are also useful when doing major changes to the database.

SELECT...FOR UPDATE

Perform a **SELECT** and then sets an exclusive lock on the matching tuples.

Can also set shared locks:

- Postgres: **FOR SHARE**
- MySQL: **LOCK IN SHARE MODE**

Table 13.3. Conflicting Row-Level Locks

Requested Lock Mode	Current Lock Mode			
	FOR KEY SHARE	FOR SHARE	FOR NO KEY UPDATE	FOR UPDATE
FOR KEY SHARE				X
FOR SHARE			X	X
FOR NO KEY UPDATE		X	X	X
FOR UPDATE	X	X	X	X

```
SELECT * FROM <table>
WHERE <qualification> FOR UPDATE;
```

SELECT...SKIP LOCKED

Perform a **SELECT** and automatically ignore any tuples that are already locked in an incompatible mode.

→ Useful for maintaining queues inside of a DBMS.

```
SELECT * FROM <table>
WHERE <qualification> SKIP LOCKED;
```

Conclusion

2PL is used in almost every DBMS.

Automatically generates correct interleaving:

- Locks + protocol (2PL, SS2PL ...)
- Deadlock detection + handling
- Deadlock prevention

Many more things not discussed...

- Nested Transactions
- Savepoints

NEXT CLASS


Timestamp Ordering Concurrency Control
Isolation Levels

Project #3 – Query Execution

You will add support for optimizing and executing queries in BusTub.

BusTub supports (basic) SQL with a rule-based optimizer for converting AST into physical plans.

Project #3 – Query Execution



PROJECT #3 – TASKS

Plan Node Executors

- Access Methods: Sequential Scan, Index Scan
- Modifications: Insert, Delete, Update
- Joins: Nested-Loop, Index Nested-Loop Hash Join
- Miscellaneous: External Merge-Sort, Limit

Optimizer Rule:

- Convert Nested Loops to Hash Join
- Convert Sequential Scan to Index Scan

PROJECT #3 - LEADERBOARD

The leaderboard requires you to add additional rules to the optimizer to generate query plans.

→ It will be impossible to get a top ranking by just having the fastest implementations in Project #1 + Project #2.

Tasks:

- Column Pruning
- More Aggressive Predicate Pushdown
- Bloom Filter for Hash Join

DEVELOPMENT HINTS

Implement the **Insert** and **Sequential Scan** executors first so that you can populate tables and read from it.

You do not need to worry about transactions.

The aggregation hash table does not need to be backed by your buffer pool (i.e., use STL)

Gradescope is for meant for grading, not debugging. Write your own local tests.

THINGS TO NOTE

Do **not** change any file other than the ones that you submit to Gradescope.

Make sure you pull in the latest changes from the BusTub main branch.

Come to TA office hours.

Compare against our solution in your browser!